

A decorative network diagram in the top-left corner of the slide. It features a complex web of interconnected nodes and edges. The nodes are represented by small circles, some of which are solid blue, some are solid grey, and some are hollow with a blue outline. The edges are thin grey lines connecting the nodes. The overall structure is a dense, interconnected mesh.

# **BST 261: Data Science II**

## **Lecture 2**

**MLPs, Backpropagation, and coding neural  
networks in Python**

**Heather Mattie**  
**Harvard T.H. Chan School of Public Health**  
**Spring 2 2021**

A decorative network diagram in the bottom-right corner of the slide. It features a complex web of interconnected nodes and edges. The nodes are represented by small circles, some of which are solid blue, some are solid grey, and some are hollow with a blue outline. The edges are thin grey lines connecting the nodes. The overall structure is a dense, interconnected mesh.

# Recipe of the Day


## Blueberry scones with lemon glaze



© Sally's Baking Addiction

# Deep learning glossaries

1. [Google](#)
2. [WildML](#)






# Perceptrons

# Perceptrons


- Let's put this all together
- Our first network will be a single neuron that will learn a simple function



$X_1$	$X_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	1

# Perceptrons

- How do we make a prediction for each observation?



Assume the following values:

$w_1$	$w_2$	$b$
1	-1	-0.5

Observations

$X_1$	$X_2$	$y$
0	0	0
0	1	1
1	0	1
1	1	1

# Predictions

- For the first observation,  $X_1 = 0, X_2 = 0, y = 0$
- First compute the weighted sum:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$h = 1 * 0 + -1 * 0 + (-0.5)$$

$$h = -0.5$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Observations

x1	x2	y
0	0	0
0	1	1
1	0	1
1	1	1

# Predictions

⊙ For the first observation,  $X_1 = 0, X_2 = 0, y = 0$

⊙ First compute the weighted sum:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$h = 1 * 0 + -1 * 0 + (-0.5)$$

$$h = -0.5$$

Transform to a probability:

$$p = \frac{1}{1 + \exp(-h)}$$

$$p = \frac{1}{1 + \exp(-0.5)}$$

$$p = 0.38$$

Assume the following values:

w1	w2	b
1	-1	-0.5

\*Note we are doing binary classification so we have to use the sigmoid activation function to calculate p



# Predictions

⊙ For the first observation,  $X_1 = 0, X_2 = 0, y = 0$

⊙ First compute the weighted sum:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$h = 1 * 0 + -1 * 0 + (-0.5)$$

$$h = -0.5$$

Transform to a probability:

$$p = \frac{1}{1 + \exp(-h)}$$

$$p = \frac{1}{1 + \exp(-0.5)}$$

$$p = 0.38$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Round to get prediction:

$$\hat{y} = \text{round}(p)$$

$$\hat{y} = 0$$

# Predictions

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$\hat{y} = \text{round}(p)$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Complete the table:

X1	X2	y	h	p	$\hat{y}$
0	0	0	-0.5	0.38	0
0	1	1			
1	0	1			
1	1	1			

# Predictions

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

$$\hat{y} = \text{round}(p)$$

Assume the following values:

w1	w2	b
1	-1	-0.5

Complete the table:

X1	X2	y	h	p	$\hat{y}$
0	0	0	-0.5	0.38	0
0	1	1	-1.5	0.18	0
1	0	1	0.5	0.38	0
1	1	1	-0.5	0.38	0

# Performance

- ◎ Our network isn't so great
- ◎ How do we make it better?
- ◎ What does *better* mean?
  - Need to define a measure of performance
  - There are many ways
- ◎ Let's begin with squared error:  $(y - p)^2$
- ◎ We need to find values for  $w_1, w_2, b$  that make this error as small as possible.
- ◎ We need to **learn** values for  $w_1, w_2, b$  such that the difference between the predicted and actual values is as small as possible.

# Learning From Data


How do we find the best values for  $w_1, w_2, b$  ?

## How do we find the best values for $w_1, w_2, b$ ?



# Learning From Data

How do we find





# Learning From Data

- ◎ Recall that the derivative of a function tells you how it is changing at any given location.
  - If the derivative is positive, it means it's going up
  - If the derivative is negative, it means it's going down
- ◎ Strategy:
  - Start with initial values for  $w_1, w_2, b$
  - Take partial derivatives of the loss function with respect to  $w_1, w_2, b$
  - Subtract the derivative (also called the **gradient**) from each
  - This is known as **gradient descent**


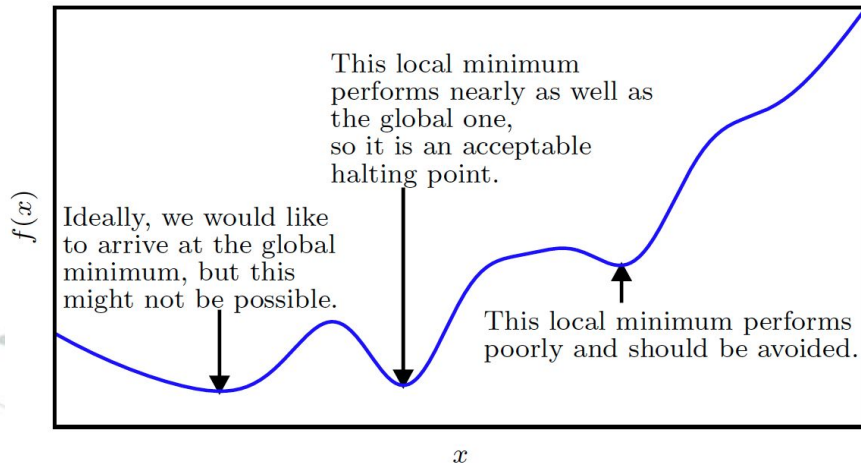


# Gradient-Based Optimization



# Gradient-Based Optimization

- ⊙ A point that obtains the absolute lowest value of  $f(x)$  is a global minimum
- ⊙ There may be one global minimum or multiple global minima
- ⊙ It is also possible for there to be local minima that are not globally optimal
- ⊙ It is common in many settings to settle for a value  $f$  that is very low but not necessarily minimal



# Gradient-Based Optimization

- ◎ To minimize  $f$ , we would like to find the direction in which  $f$  decreases the fastest
- ◎ It can be shown that the gradient points directly uphill and the negative gradient directly downhill
- ◎ We can therefore decrease  $f$  by moving in the direction of the negative gradient
- ◎ For example, for a weight  $w_i$

$$w_i^{\text{new}} = w_i^{\text{old}} - \eta g$$

where  $\eta$  is the **learning rate** (how fast you want to move down the gradient), and  $g$  is the gradient

# The Backpropagation Algorithm

## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton†  
& Ronald J. Williams\*

\* Institute for Cognitive Science, C-015, University of California,  
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,  
Pittsburgh, Philadelphia 15213, USA

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure<sup>1</sup>.



# The Backpropagation Algorithm

- Our perceptron performs the following computations:

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

- We want to minimize this quantity:

$$l = (y - p)^2$$

- We'll compute the gradients for each parameter by “backpropagating” errors through each component of the network

# The Backpropagation Algorithm

For  $w_1$  we need to compute

$$\frac{\partial l}{\partial w_1}$$

To get there we will use the chain rule

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

This is “backprop”

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$l = (y - p)^2$$

# The Backpropagation Algorithm

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\boxed{\frac{\partial l}{\partial p}} =$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$\boxed{l = (y - p)^2}$$

# The Backpropagation Algorithm

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\boxed{\frac{\partial l}{\partial p}} = 2 * (p - y)$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$\boxed{l = (y - p)^2}$$



# The Backpropagation Algorithm

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\boxed{\frac{\partial l}{\partial p}} = 2 * (p - y)$$

$$\boxed{\frac{\partial p}{\partial h}} =$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$\boxed{p = \frac{1}{1 + \exp(-h)}}$$

Loss

$$\boxed{l = (y - p)^2}$$

# The Backpropagation Algorithm

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\boxed{\frac{\partial l}{\partial p}} = 2 * (p - y)$$

$$\boxed{\frac{\partial p}{\partial h}} = p * (1 - p)$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$\boxed{p = \frac{1}{1 + \exp(-h)}}$$

Loss

$$\boxed{l = (y - p)^2}$$

# The Backpropagation Algorithm

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2 * (p - y)$$

$$\frac{\partial p}{\partial h} = p * (1 - p)$$

$$\frac{\partial h}{\partial w_1} =$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$l = (y - p)^2$$

# The Backpropagation Algorithm

$$\frac{\partial l}{\partial w_1} = \frac{\partial l}{\partial p} * \frac{\partial p}{\partial h} * \frac{\partial h}{\partial w_1}$$

$$\frac{\partial l}{\partial p} = 2 * (p - y)$$

$$\frac{\partial p}{\partial h} = p * (1 - p)$$

$$\frac{\partial h}{\partial w_1} = X_1$$

Computations

$$h = w_1 * X_1 + w_2 * X_2 + b$$

$$p = \frac{1}{1 + \exp(-h)}$$

Loss

$$l = (y - p)^2$$

Putting it all together:

$$\frac{\partial l}{\partial w_1} = 2 * (p - y) * p * (1 - p) * X_1$$

# Gradient Descent with Backprop

For some number of iterations:

1. Compute the gradient for  $w_1$
2. Update  $w_i^{\text{new}} = w_i^{\text{old}} - \eta g$
3. Repeat until “convergence”

Do this for each weight and bias term.



# Multilayer Perceptrons

# Perceptron → MLP


We can turn our perceptron model into a multilayer perceptron

- ⊙ Instead of just one linear combination, we are going to take several, each with a different set of weights
- ⊙ Each linear combination will be followed by a nonlinear activation
- ⊙ Each of these nonlinear features will be fed into the logistic regression classifier (binary classifier)
- ⊙ All of the weights are learned end-to-end via SGD

MLPs learn a set of nonlinear features directly from data - “feature engineering” is the hallmark of deep learning approaches

# Multilayer Perceptrons (MLPs)

Suppose we have the following MLP with 1 hidden layer that has 3 hidden units:



Each neuron in the hidden layer is going to do exactly the same thing as before.




# Multilayer Perceptrons (MLPs)

Computations:

$$h_j = \phi(w_{1j} * X_1 + w_{2j} * X_2 + b_j)$$

$$o = b_o + \sum_{j=1}^3 w_{oj} * h_j$$

$$p = \frac{1}{1 + \exp(-o)}$$



\*If we use a sigmoid activation function

Output layer weight derivatives


$$\begin{aligned} \frac{\partial l}{\partial w_{oj}} &= \frac{\partial l}{\partial p} * \frac{\partial p}{\partial o} * \frac{\partial o}{\partial w_{oj}} \\ &= (p - y) * p * (1 - p) * h_j \end{aligned}$$

Hidden layer weight derivatives

$$\begin{aligned} \frac{\partial l}{\partial w_{1j}} &= \frac{\partial l}{\partial p} * \frac{\partial p}{\partial o} * \frac{\partial o}{\partial h} * \frac{\partial h}{\partial w_{1j}} \\ &= (p - y) * p * (1 - p) * h_j * (1 - h_j) * X_1 \end{aligned}$$

# Matrix Notation

Sum notation starts to get unwieldy quickly. We can use matrix notation to represent each calculation in a more concise way.




$$X = \begin{bmatrix} X_1 \\ X_2 \end{bmatrix} \quad W = \begin{bmatrix} w_{1,1} & w_{2,1} & w_{3,1} \\ w_{1,2} & w_{2,2} & w_{3,2} \end{bmatrix} \quad B = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

$$Z = W^T X + B \quad H = \begin{bmatrix} h_1 \\ h_2 \\ h_3 \end{bmatrix} = \phi(Z)$$

# Notation


As the number of layers grows, the number of matrices grows and we have to add a superscript to denote the layer. We also have to add a superscript to denote which training example we are referencing.

Example notation for 1 weight in 1 hidden layer for 1 training example:




# MLP Terminology

Forward pass = computing probability from input



# MLP Terminology


Forward pass = computing probability from input



Backward pass = computing derivatives from the output

# MLP Terminology

Forward pass = computing probability from input




Hidden layers are also called “dense” layers or “fully connected” layers

Backward pass = computing derivatives from the output

# MLPs

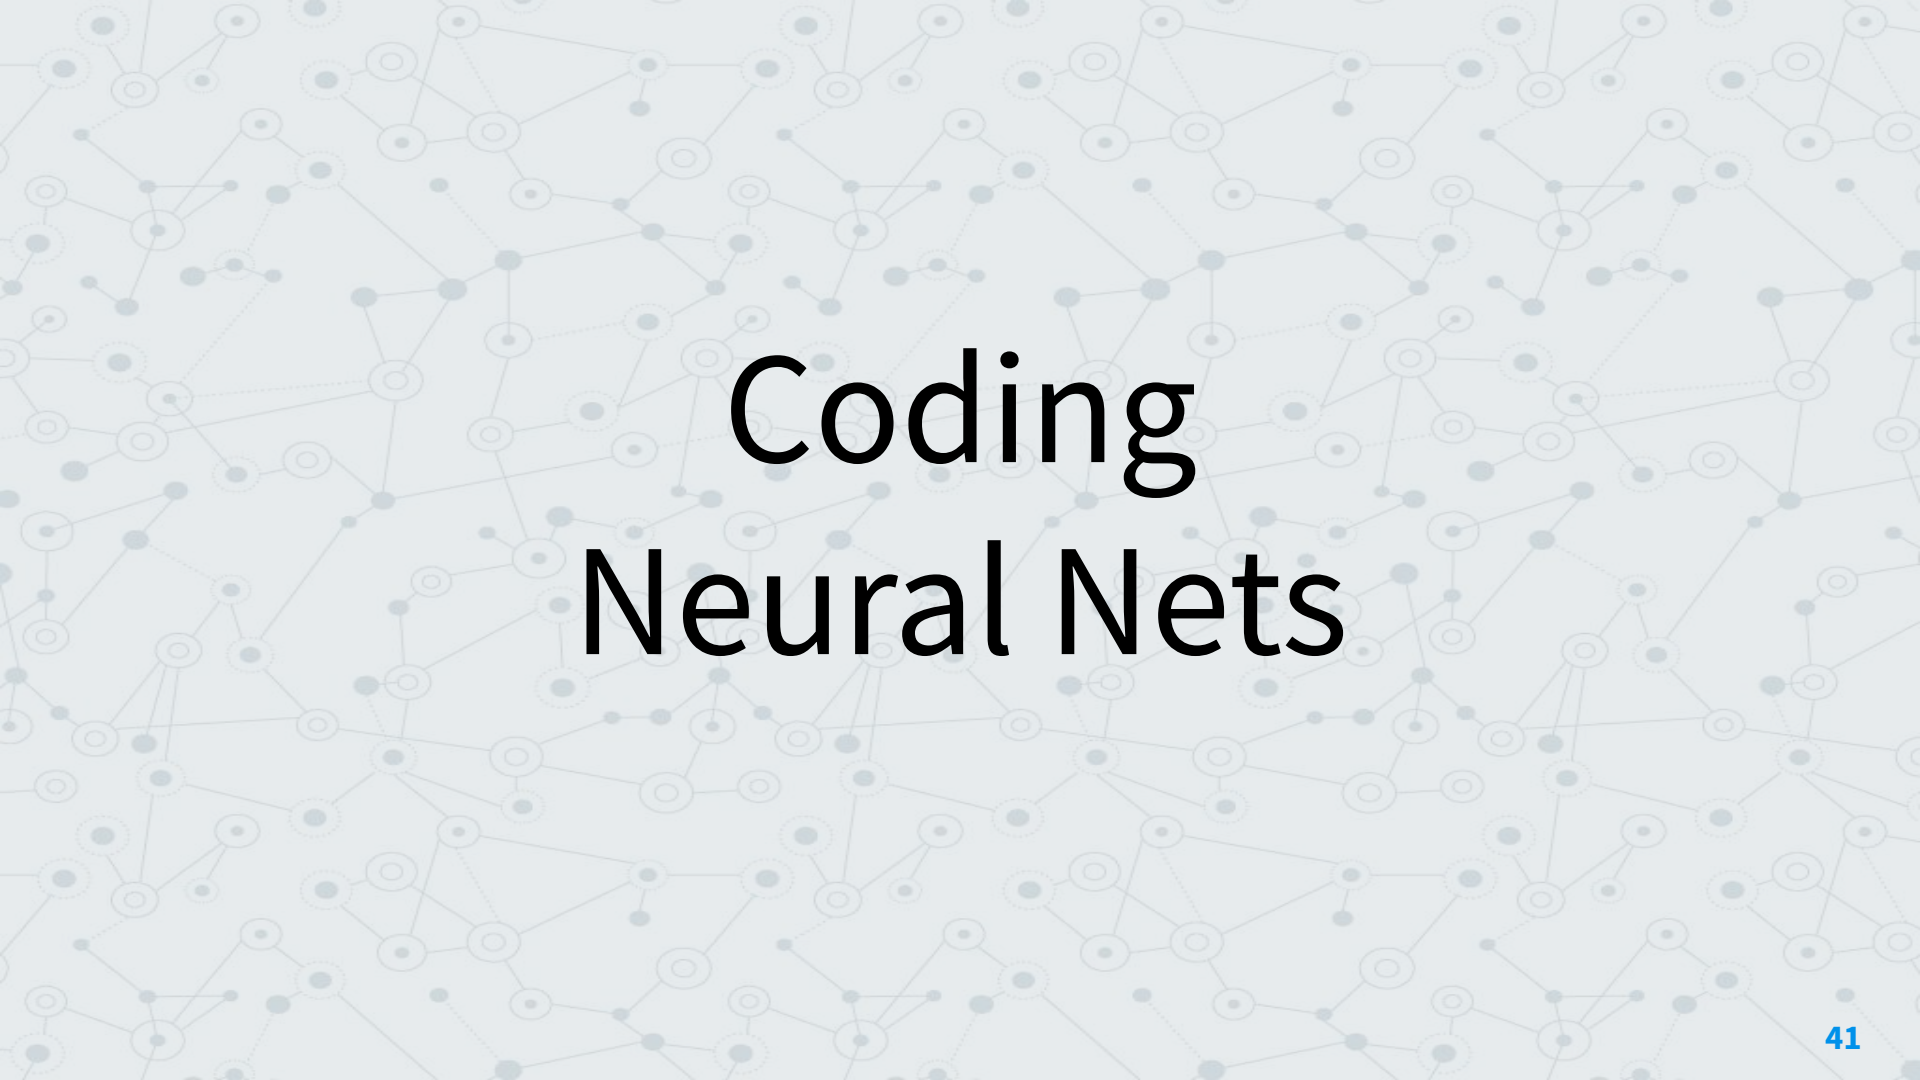
Increasing the number of layers increases the flexibility of the model - but run the risk of overfitting



# Conclusions

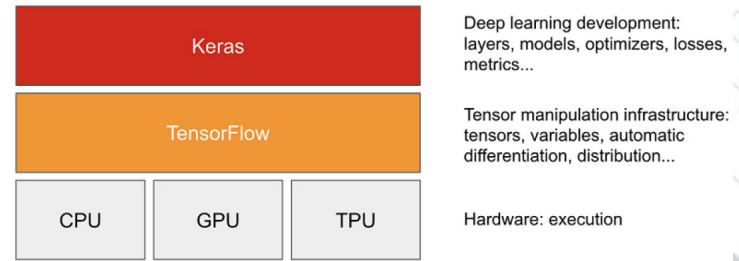
- ◎ Backprop, perceptrons, and MLPs are the building blocks of neural nets
- ◎ You'll get a chance to demonstrate your mastery in Problem Set 1
- ◎ We will use these concepts for the rest of the semester





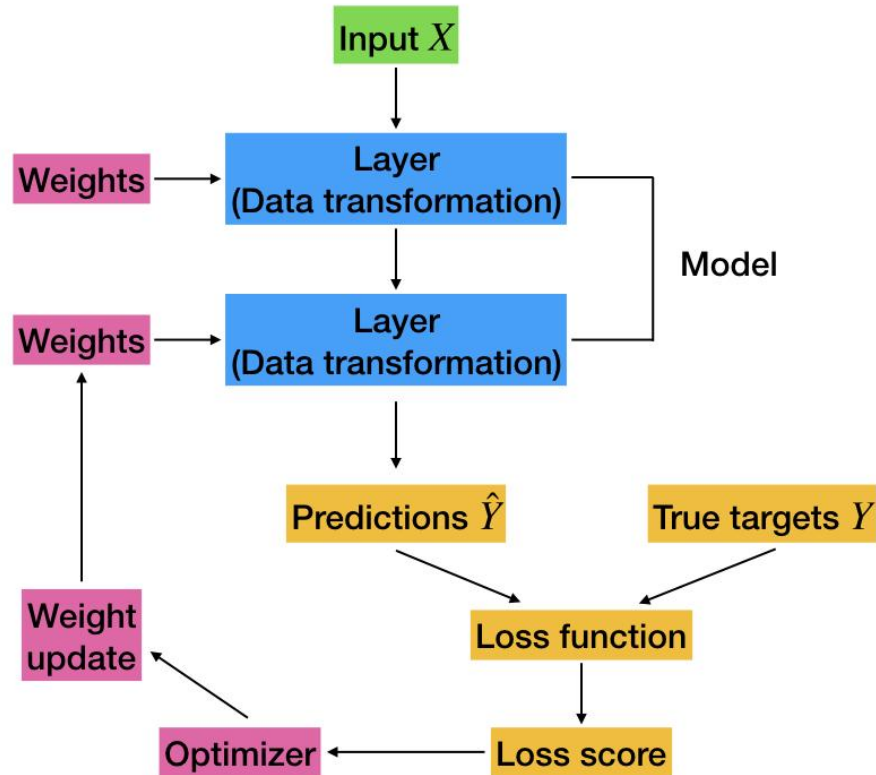
# Coding Neural Nets

# Keras and Tensorflow



- ◎ Keras is a model-level library that provides high-level building blocks for developing deep learning models
- ◎ It doesn't handle low-level operations like matrix and tensor (n-dimensional matrix) multiplication and differentiation
  - It uses TensorFlow or Theano or CNTK (Microsoft Cognitive Toolkit) backends for this
  - We will be using TensorFlow
    - ◎ It is the most widely adopted, scalable and production ready
- ◎ Keras can run on both CPUs and GPUs
  - When running on CPUs, uses Eigen for tensor operations
  - When running on GPUs, uses the NVIDIA CUDA Deep Neural Network library (cuDNN)

# Neural Network Workflow



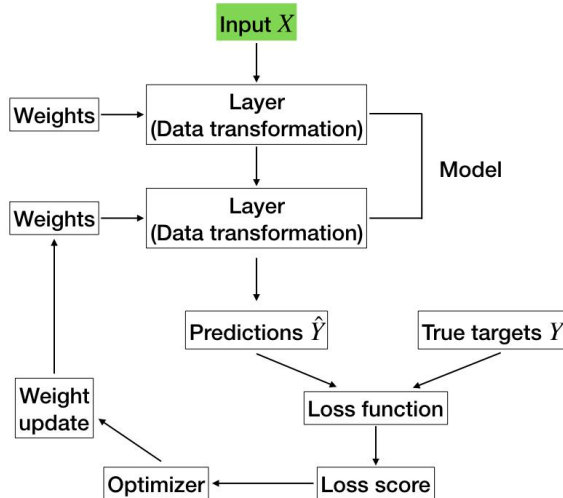
# Generic Feedforward Network

Elements needed:


1. Necessary libraries
2. Dataset split into training and test sets (validation as well if you have enough data)
3. `models.Sequential()` : defines a linear, or sequential architecture made up of a set of layers that will stack to create the network
4. `layers.Dense()` : specifies a fully connected layer
5. `model.compile(optimizer, loss, metrics)` : specifies how to execute the training of the network
6. `model.fit(train_data, train_labels, epochs, batch_size)` : fits the neural net using the training data, runs for a specified number of iterations (epochs) using batch\_size number of training examples at a time

# Generic Feedforward Network

```
1 # Import needed packages (not an exhaustive list)
2 import tensorflow as tf
3 from tensorflow import keras
4 from tensorflow.keras import layers
5
6 # Load data (will most likely be more complicated)
7 (x_train, y_train), (x_test, y_test) = load_data()
8
9 # Define model architecture
10 model = keras.Sequential([
11     # Layer 1 (Hidden layer, fully connected)
12     layers.Dense(c, activation='activation function'),
13     # Layer 2 (Output layer, fully connected)
14     layers.Dense(d, activation='output activation function')
15 ])
16
17 # Define how to execute training
18 model.compile(optimizer='optimizing algorithm',
19               loss='loss function',
20               metrics=['performance metric'])
21
22 # Train the network
23 model.fit(x_train, y_train, epochs = e, batch_size = b)
```




```
1 # Import needed packages (not an exhaustive list)
2 import tensorflow as tf
3 from tensorflow import keras
4 from tensorflow.keras import layers
5
6 # Load data (will most likely be more complicated)
7 (x_train, y_train), (x_test, y_test) = load_data()
8
9 # Define model architecture
10 model = keras.Sequential([
11     # Layer 1 (Hidden layer, fully connected)
12     layers.Dense(c, activation='activation function'),
13     # Layer 2 (Output layer, fully connected)
14     layers.Dense(d, activation='output activation function')
15 ])
16
17 # Define how to execute training
18 model.compile(optimizer='optimizing algorithm',
19               loss='loss function',
20               metrics=['performance metric'])
21
22 # Train the network
23 model.fit(x_train, y_train, epochs = e, batch_size = b)
```



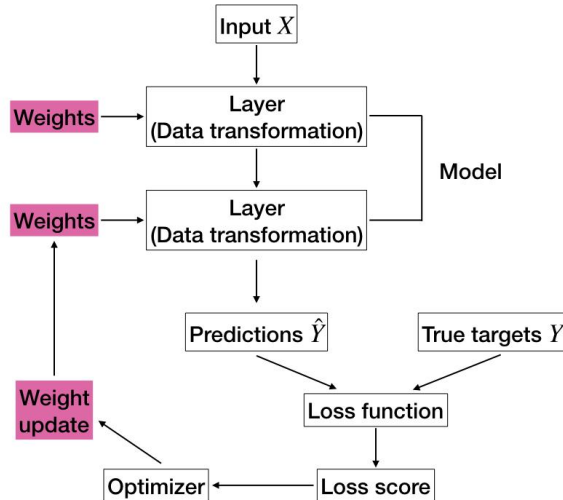
```
1 # Import needed packages (not an exhaustive list)
2 import tensorflow as tf
3 from tensorflow import keras
4 from tensorflow.keras import layers
5
6 # Load data (will most likely be more complicated)
7 (x_train, y_train), (x_test, y_test) = load_data()
8
9 # Define model architecture
10 model = keras.Sequential([
11     # Layer 1 (Hidden layer, fully connected)
12     layers.Dense(c, activation='activation function'),
13     # Layer 2 (Output layer, fully connected)
14     layers.Dense(d, activation='output activation function')
15 ])
16
17 # Define how to execute training
18 model.compile(optimizer='optimizing algorithm',
19               loss='loss function',
20               metrics=['performance metric'])
21
22 # Train the network
23 model.fit(x_train, y_train, epochs = e, batch_size = b)
```





```
1 # Import needed packages (not an exhaustive list)
2 import tensorflow as tf
3 from tensorflow import keras
4 from tensorflow.keras import layers
5
6 # Load data (will most likely be more complicated)
7 (x_train, y_train), (x_test, y_test) = load_data()
8
9 # Define model architecture
10 model = keras.Sequential([
11     # Layer 1 (Hidden layer, fully connected)
12     layers.Dense(c, activation='activation function'),
13     # Layer 2 (Output layer, fully connected)
14     layers.Dense(d, activation='output activation function')
15 ])
16
17 # Define how to execute training
18 model.compile(optimizer='optimizing algorithm',
19               loss='loss function',
20               metrics=['performance metric'])
21
22 # Train the network
23 model.fit(x_train, y_train, epochs = e, batch_size = b)
```





```
1 # Import needed packages (not an exhaustive list)
2 import tensorflow as tf
3 from tensorflow import keras
4 from tensorflow.keras import layers
5
6 # Load data (will most likely be more complicated)
7 (x_train, y_train), (x_test, y_test) = load_data()
8
9 # Define model architecture
10 model = keras.Sequential([
11     # Layer 1 (Hidden layer, fully connected)
12     layers.Dense(c, activation='activation function'),
13     # Layer 2 (Output layer, fully connected)
14     layers.Dense(d, activation='output activation function')
15 ])
16
17 # Define how to execute training
18 model.compile(optimizer='optimizing algorithm',
19               loss='loss function',
20               metrics=['performance metric'])
21
22 # Train the network
23 model.fit(x_train, y_train, epochs = e, batch_size = b)
```

# Generic Feedforward Network

**train\_data**: training examples (matrix of feature vectors;  $\mathbf{X}_{\text{train}}$ )

**train\_labels**: training labels ( $\mathbf{y}_{\text{train}}$ )

**test\_data**: test examples used to measure performance of network ( $\mathbf{X}_{\text{test}}$ )

**test\_labels**: test set labels ( $\mathbf{y}_{\text{test}}$ )

Optimizing algorithms: rmsprop, sgd, adagrad, adam, etc.

Loss function options: mse, mae, categorical\_crossentropy, etc.

Performance measure options: accuracy, mae, etc.

Here:

**c** = the number of hidden units (neurons) in a hidden layer

**d** = the number of units (neurons) in the output layer

**e** = the number of epochs (iterations) over entire training data set

**b** = the batch size (how many training examples to optimize at once)