

A decorative network diagram in the top-left corner, featuring a complex web of interconnected nodes and lines. Some nodes are highlighted with blue circles or dots.

BST 261: Data Science II

Lecture 5

**Convolutional Neural Networks (CNNs):
Data Augmentation and Pretrained Bases**

Heather Mattie
Harvard T.H. Chan School of Public Health
Spring 2 2021

A decorative network diagram in the bottom-right corner, featuring a complex web of interconnected nodes and lines. Some nodes are highlighted with blue circles or dots.

Recipe of the Day!

Iced vanilla and caramel profiteroles






CNNs in Python

Classifying Skin Lesions

- ◎ We'll be using data from the [International Skin Imaging Collaboration: Melanoma Project](#)
- ◎ We'll be classifying images as malignant or benign
- ◎ The overarching goal of the ISIC Melanoma Project is to support efforts to reduce melanoma-related deaths and unnecessary biopsies by improving the accuracy and efficiency of melanoma early detection

Malignant



Malignant



Benign



Benign



Classifying Skin Lesions

- ◎ This archive contains 23k images of classified skin lesions. It contains both malignant and benign examples
 - We'll be using a fraction of this
- ◎ Each example contains the image of the lesion, meta data regarding the lesion (including classification and segmentation) and meta data regarding the patient
- ◎ The data can be viewed in [this link](#) (in the gallery section)
- ◎ It can be downloaded through the site or by using [this repository](#)

Classification Skin Lesions

- ◎ The subsample of the data is available in a [Google Drive](#) folder
- ◎ You can access it with this [code and notebook](#)
- ◎ You can also download the images to your machine if you would like
 - There are zip files available on canvas
- ◎ We'll start with creating a simple CNN

```

1 # Define model
2 model = keras.Sequential([
3     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
4     layers.MaxPooling2D((2, 2)),
5
6     layers.Conv2D(64, (3, 3), activation='relu'),
7     layers.MaxPooling2D((2, 2)),
8
9     layers.Conv2D(128, (3, 3), activation='relu'),
10    layers.MaxPooling2D((2, 2)),
11
12    layers.Conv2D(128, (3, 3), activation='relu'),
13    layers.MaxPooling2D((2, 2)),
14
15    layers.Flatten(),
16
17    layers.Dense(512, activation='relu'),
18
19    layers.Dense(1, activation='sigmoid')
20 ])

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513
=====		


Total params: 3,453,121
 Trainable params: 3,453,121
 Non-trainable params: 0


```

1 # Define model
2 model = keras.Sequential([
3     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
4     layers.MaxPooling2D((2, 2)),
5
6     layers.Conv2D(64, (3, 3), activation='relu'),
7     layers.MaxPooling2D((2, 2)),
8
9     layers.Conv2D(128, (3, 3), activation='relu'),
10    layers.MaxPooling2D((2, 2)),
11
12    layers.Conv2D(128, (3, 3), activation='relu'),
13    layers.MaxPooling2D((2, 2)),
14
15    layers.Flatten(),
16
17    layers.Dense(512, activation='relu'),
18
19    layers.Dense(1, activation='sigmoid')
20 ])

```

Convolution and pooling layers - notice how the output size decreases with each layer. Remember what applying a filter, padding, and/or strides does to an input.



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513
=====		
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		


```

1 # Define model
2 model = keras.Sequential([
3     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
4     layers.MaxPooling2D((2, 2)),
5
6     layers.Conv2D(64, (3, 3), activation='relu'),
7     layers.MaxPooling2D((2, 2)),
8
9     layers.Conv2D(128, (3, 3), activation='relu'),
10    layers.MaxPooling2D((2, 2)),
11
12    layers.Conv2D(128, (3, 3), activation='relu'),
13    layers.MaxPooling2D((2, 2)),
14
15    layers.Flatten(),
16
17    layers.Dense(512, activation='relu'),
18
19    layers.Dense(1, activation='sigmoid')
20 ])

```

We finish the network with 1 hidden dense layer and 1 output layer. Note that most of the parameters in the model come from the hidden dense layer and not the convolution or pooling layers.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513

Total params: 3,453,121
Trainable params: 3,453,121
Non-trainable params: 0

```

1 # Define model
2 model = keras.Sequential([
3     layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
4     layers.MaxPooling2D((2, 2)),
5
6     layers.Conv2D(64, (3, 3), activation='relu'),
7     layers.MaxPooling2D((2, 2)),
8
9     layers.Conv2D(128, (3, 3), activation='relu'),
10    layers.MaxPooling2D((2, 2)),
11
12    layers.Conv2D(128, (3, 3), activation='relu'),
13    layers.MaxPooling2D((2, 2)),
14
15    layers.Flatten(),
16
17    layers.Dense(512, activation='relu'),
18
19    layers.Dense(1, activation='sigmoid')
20 ])

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_1 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_2 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_3 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten (Flatten)	(None, 6272)	0
dense (Dense)	(None, 512)	3211776
dense_1 (Dense)	(None, 1)	513

Total params: 3,453,121
 Trainable params: 3,453,121
 Non-trainable params: 0

Total # of parameters that need to be learned by the network

```
1 from keras.preprocessing.image import ImageDataGenerator
2
3 # All images will be rescaled by 1./255
4 train_datagen = ImageDataGenerator(rescale=1./255)
5 test_datagen = ImageDataGenerator(rescale=1./255)
6
7 train_generator = train_datagen.flow_from_directory(
8     # This is the target directory
9     train_dir,
10    # All images will be resized to 150x150
11    target_size = (150, 150),
12    batch_size = 20,
13    # Since we use binary_crossentropy loss, we need binary labels
14    class_mode = 'binary')
15
16 validation_generator = test_datagen.flow_from_directory(
17     validation_dir,
18     target_size = (150, 150),
19     batch_size = 20,
20     class_mode = 'binary')
```

We first scale the data to get values between 0 and 1.

Then we transform the images to be 150x150 pixels in size (this is arbitrary), declare a batch size of 20 (this is also arbitrary), and declare the class mode (i.e. the type of classification we want to do)



```
1 for data_batch, labels_batch in train_generator:
2     print('data batch shape:', data_batch.shape)
3     print('labels batch shape:', labels_batch.shape)
4     break
```

```
data batch shape: (20, 150, 150, 3)
labels batch shape: (20,)
```

We can see the shape of the images: they are in batches of 20, with each image being represented by 3, 150x150 tensors: one for R, one for G and one for B color “channels”.

```
1 history = model.fit(
2     train_generator,
3     steps_per_epoch = 81, # ceil(1609/20)
4     epochs = 30,
5     validation_data = validation_generator,
6     validation_steps = 22) # ceil(426/20)
```

The number of training examples divided by the batch size, i.e. how many batches we need to go through until the model sees all training data. For both the training and validation sets.



Data Augmentation

- ◎ As we have seen, overfitting is caused by having too few training examples to learn from
- ◎ Data augmentation generates more training data from existing training examples by **augmenting** the samples via a number of random transformations
- ◎ These transformations should yield believable images

Data Augmentation

- ◎ Types of augmentation:
 - Rotation
 - Horizontal/vertical flip
 - Random crops/scales
 - ◎ Zoom
 - ◎ Width or height shifts
 - Shearing
 - Brightness, contrast, saturation
 - Lens distortions

Types of data augmentation

1. Rotations



Types of data augmentation

2. Horizontal/Vertical Flips




Types of data augmentation

3. Random crops/scales



Types of data augmentation

4. Shearing



Types of data augmentation

5. Brightness, contrast, saturation



Types of data augmentation

6. Lens distortions



Types of data augmentation

7. Combinations of the above



Data Augmentation

- ◎ If you train a network using data augmentation, it will never see the same input twice, but the inputs will still be heavily correlated
 - You're remixing known information, not producing new information
- ◎ May not completely escape overfitting due to this correlation
- ◎ Adding dropout can also help

Data Augmentation in Keras





```
1 from keras.preprocessing.image import ImageDataGenerator
2 datagen = ImageDataGenerator(
3     rotation_range = 40,
4     width_shift_range = 0.2,
5     height_shift_range = 0.2,
6     shear_range = 0.2,
7     zoom_range = 0.2,
8     horizontal_flip = True,
9     fill_mode = 'nearest')
```

You can create your own data generator with any specifications you'd like. The values chosen here are arbitrary.

You can check out the [Keras documentation](#) to see all of the available options and values each type of augmentation type can take.

Note that only your training data should be augmented - not the test or validation sets. The point of augmentation is to “increase” your training set size.

Data Augmentation in Keras




Back to the notebook

[Skin lesions with data augmentation](#)

Data Augmentation in Keras

Without augmentation



With augmentation

